

**Московский государственный технический  
университет им. Н.Э. Баумана.**

Факультет «Информатика и управление»

Кафедра ИУ5. Курс «Технологии машинного обучения»

Отчет по лабораторной работе №2

Выполнила:		Проверил:
студентка группы ИУ5-61		преподаватель каф. ИУ5
Абросимова Надежда		Гапанюк Ю.Е.
Подпись и дата:		Подпись и дата:

г. Москва, 2019 г.

## Задание

### Часть 1.

Выполните первое демонстрационное задание "demo assignment" под названием "Exploratory data analysis with Pandas" со страницы курса <https://mlcourse.ai/assignments>

Условие задания -

[https://nbviewer.jupyter.org/github/Yorko/mlcourse\\_open/blob/master/jupyter\\_english/assignments\\_demo/assignment01\\_pandas\\_uci\\_adult.ipynb?flush\\_cache=true](https://nbviewer.jupyter.org/github/Yorko/mlcourse_open/blob/master/jupyter_english/assignments_demo/assignment01_pandas_uci_adult.ipynb?flush_cache=true)

Набор данных можно скачать здесь - <https://archive.ics.uci.edu/ml/datasets/Adult>

Пример решения задания - <https://www.kaggle.com/kashnitsky/a1-demo-pandas-and-uci-adult-dataset-solution>

### Часть 2.

Выполните следующие запросы с использованием двух различных библиотек - Pandas и PandaSQL:

один произвольный запрос на соединение двух наборов данных  
один произвольный запрос на группировку набора данных с использованием функций агрегирования  
Сравните время выполнения каждого запроса в Pandas и PandaSQL.

## Текст программы

*#Часть 1*

```
import numpy as np
import pandas as pd
```

```
data = pd.read_csv('adult.data.csv')
```

```
data.head()
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40

*#1. How many men and women (sex feature) are represented in this dataset?*

```
data["sex"].value_counts()
```

```
Out[4]:
```

```
Male      21790
Female    10771
```

Name: sex, dtype: int64

In [5]:

*#2. What is the average age (age feature) of women?*

```
data.loc[data['sex'] == 'Female', 'age'].mean()
```

Out[5]:

36.85823043357163

In [6]:

*#3. What is the proportion of German citizens (native-country feature)?*

```
print("{0:%}".format(data[data["native-country"] == "Germany"]  
                        .shape[0] / data.shape[0]))
```

0.420749%

In [7]:

*#4-5. What are the mean and standard deviation of age for those who earn more than 50K per year (salary feature)*

*#and those who earn less than 50K per year?*

```
ages1 = data[data["salary"] == "<=50K"]["age"]  
ages2 = data[data["salary"] == ">50K"]["age"]  
print("<=50K: = {0} ± {1} years".format(ages1.mean(), ages1.std()))  
print(">50K: = {0} ± {1} years".format(ages2.mean(), ages2.std()))  
<=50K: = 36.78373786407767 ± 14.020088490824813 years  
>50K: = 44.24984058155847 ± 10.51902771985177 years
```

In [8]:

*#6. Is it true that people who earn more than 50K have at least high school education?*

*#(education - Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)*

```
high_educations = set(["Bachelors", "Prof-school", "Assoc-acdm",  
                       "Assoc-voc", "Masters", "Doctorate"])
```

```
def high_educated(e):  
    return e in high_educations
```

```
data[data["salary"] == ">50K"]["education"].map(high_educated).all()
```

Out[8]:

False

In [9]:

*#7. Display statistics of age for each race (race feature) and each gender. Use groupby() and describe().*

*#Find the maximum age of men of Amer-Indian-Eskimo race.*

```
for (race, sex), sub_df in data.groupby(['race', 'sex']):  
    print("Race: {0}, sex: {1}".format(race, sex))  
    print(sub_df['age'].describe())
```

Race: Amer-Indian-Eskimo, sex: Female

count	119.000000
mean	37.117647
std	13.114991
min	17.000000
25%	27.000000
50%	36.000000
75%	46.000000

```
max      80.000000
Name: age, dtype: float64
Race: Amer-Indian-Eskimo, sex: Male
count    192.000000
mean     37.208333
std      12.049563
min      17.000000
25%      28.000000
50%      35.000000
75%      45.000000
max      82.000000
Name: age, dtype: float64
Race: Asian-Pac-Islander, sex: Female
count    346.000000
mean     35.089595
std      12.300845
min      17.000000
25%      25.000000
50%      33.000000
75%      43.750000
max      75.000000
Name: age, dtype: float64
Race: Asian-Pac-Islander, sex: Male
count    693.000000
mean     39.073593
std      12.883944
min      18.000000
25%      29.000000
50%      37.000000
75%      46.000000
max      90.000000
Name: age, dtype: float64
Race: Black, sex: Female
count    1555.000000
mean     37.854019
std      12.637197
min      17.000000
25%      28.000000
50%      37.000000
75%      46.000000
max      90.000000
Name: age, dtype: float64
Race: Black, sex: Male
count    1569.000000
mean     37.682600
std      12.882612
min      17.000000
25%      27.000000
50%      36.000000
75%      46.000000
```

```

max          90.000000
Name: age, dtype: float64
Race: Other, sex: Female
count        109.000000
mean         31.678899
std          11.631599
min          17.000000
25%          23.000000
50%          29.000000
75%          39.000000
max          74.000000
Name: age, dtype: float64
Race: Other, sex: Male
count        162.000000
mean         34.654321
std          11.355531
min          17.000000
25%          26.000000
50%          32.000000
75%          42.000000
max          77.000000
Name: age, dtype: float64
Race: White, sex: Female
count       8642.000000
mean         36.811618
std          14.329093
min          17.000000
25%          25.000000
50%          35.000000
75%          46.000000
max          90.000000
Name: age, dtype: float64
Race: White, sex: Male
count      19174.000000
mean         39.652498
std          13.436029
min          17.000000
25%          29.000000
50%          38.000000
75%          49.000000
max          90.000000
Name: age, dtype: float64

```

In [10]:

```

#8. Among whom is the proportion of those who earn a lot (>50K) greater: married or single men (marital-status feature)?
#Consider as married those who have a marital-status starting with Married (Married-civ-spouse, Married-spouse-absent or Married-AF-spouse), the rest are considered bachelors.
data.loc[(data['sex'] == 'Male') &
         (data['marital-status'].isin(['Never-married',

```

```

        'Separated',
        'Divorced',
        'Widowed'])), 'salary'].value_counts()

Out[10]:
<=50K      7552
>50K        697
Name: salary, dtype: int64

In [11]:
data.loc[(data['sex'] == 'Male') &
         (data['marital-status'].str.startswith('Married'))], 'salary'].value_counts()

Out[11]:
<=50K      7576
>50K       5965
Name: salary, dtype: int64

In [12]:
data['marital-status'].value_counts()

Out[12]:
Married-civ-spouse      14976
Never-married           10683
Divorced                 4443
Separated                1025
Widowed                  993
Married-spouse-absent    418
Married-AF-spouse        23
Name: marital-status, dtype: int64

In [13]:
#9. What is the maximum number of hours a person works per week (hours-per-week feature)?
#How many people work such a number of hours and what is the percentage of those who earn a lot among them?
max_load = data['hours-per-week'].max()
print("Max time - {0} hours./week.".format(max_load))

num_workaholics = data[data['hours-per-week'] == max_load].shape[0]
print("Total number of such hard workers {0}".format(num_workaholics))

rich_share = float(data[(data['hours-per-week'] == max_load)
                        & (data['salary'] == '>50K')].shape[0]) / num_workaholics
print("Percentage of rich among them {0}%".format(int(100 * rich_share)))
Max time - 99 hours./week.
Total number of such hard workers 85
Percentage of rich among them 29%

In [14]:
#10. Count the average time of work (hours-per-week) those who earning a little and a lot (salary) for each country
#(native-country).
pd.crosstab(data['native-country'], data['salary'],
            values=data['hours-per-week'], aggfunc=np.mean).T

```

native-country	?	Cambodia	Canada	China	Columbia	Cuba	Dominican-Republic	Ecuador	El-Salvador	England
salary										
<=50K	40.164760	41.416667	37.914634	37.381818	38.684211	37.985714	42.338235	38.041667	36.030928	40.483333
>50K	45.547945	40.000000	45.641026	38.900000	50.000000	42.440000	47.000000	48.750000	45.000000	44.533333

2 rows x 42 columns

#Часть 2

```
!pip install pandasql
from pandasql import sqldf
pysqldf = lambda q: sqldf(q, globals())
user_usage = pd.read_csv('user_usage.csv')
user_device = pd.read_csv('user_device.csv')
user_usage.head()
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id
0	21.97	4.82	1557.33	22787
1	1710.08	136.88	7267.55	22788
2	1710.08	136.88	7267.55	22789
3	94.46	35.17	519.12	22790
4	71.59	79.26	1557.33	22792

user\_usage.dtypes

Out[10]:

```
outgoing_mins_per_month    float64
outgoing_sms_per_month     float64
monthly_mb                 float64
use_id                    int64
```

dtype: object

user\_device.head()

	use_id	user_id	platform	platform_version	device	use_type_id
0	22782	26980	ios	10.2	iPhone7,2	2
1	22783	29628	android	6.0	Nexus 5	3
2	22784	28473	android	5.1	SM-G903F	1
3	22785	15200	ios	10.2	iPhone7,2	3
4	22786	28239	android	6.0	ONE E1003	1

user\_device.dtypes

Out[12]:

```
use_id          int64
user_id         int64
platform        object
platform_version float64
device          object
use_type_id     int64
```

dtype: object

In [16]:

```
user_device.merge(user_usage).head()
```

	use_id	user_id	platform	platform_version	device	use_type_id	outgoing_mins_per_month	outgoing_sms_per_month
0	22787	12921	android	4.3	GT-I9505	1	21.97	4.82
1	22788	28714	android	6.0	SM-G930F	1	1710.08	136.88
2	22789	28714	android	6.0	SM-G930F	1	1710.08	136.88
3	22790	29592	android	5.1	D2303	1	94.46	35.17
4	22792	28217	android	5.1	SM-G361F	1	71.59	79.26

```
%%timeit
```

```
user_device.merge(user_usage).head()
```

2.52 ms ± 41.1 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

In [19]:

```
pysqldf("""SELECT  u.outgoing_mins_per_month, u.outgoing_sms_per_month, u.mon
thly_mb, u.use_id, d.user_id,d.platform,
d.platform_version, d.device, d.use_type_id
FROM user_usage AS u JOIN user_device AS d
ON u.use_id = d.use_id
""").head()
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	user_id	platform	platform_version	device
0	21.97	4.82	1557.33	22787	12921	android	4.3	GT-I9505
1	1710.08	136.88	7267.55	22788	28714	android	6.0	SM-G930F
2	1710.08	136.88	7267.55	22789	28714	android	6.0	SM-G930F
3	94.46	35.17	519.12	22790	29592	android	5.1	D2303
4	71.59	79.26	1557.33	22792	28217	android	5.1	SM-G361F

```
%%timeit
```

```
pysqldf("""SELECT  u.outgoing_mins_per_month, u.outgoing_sms_per_month, u.mon
thly_mb, u.use_id, d.user_id,d.platform,
d.platform_version, d.device, d.use_type_id
FROM user_usage AS u JOIN user_device AS d
ON u.use_id = d.use_id
""")
```

13.2 ms ± 629 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

In [34]:

```
user_usage.groupby("use_id")["monthly_mb"].mean().head()
```

Out[34]:

```
use_id
22787    1557.33
22788    7267.55
22789    7267.55
22790     519.12
22792    1557.33
```



Name: monthly\_mb, dtype: float64

In [28]:

```
%%timeit
```

```
user_usage.groupby("use_id")["monthly_mb"].mean()
```

407  $\mu$ s  $\pm$  5.49  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)

In [29]:

```
pysqlldf("""SELECT use_id, AVG(monthly_mb)
```

```
FROM user_usage
```

```
GROUP BY use_id
```

```
""").head()
```

	use_id	AVG(monthly_mb)
0	22787	1557.33
1	22788	7267.55
2	22789	7267.55
3	22790	519.12
4	22792	1557.33

```
%%timeit
```

```
pysqlldf("""SELECT use_id, AVG(monthly_mb)
```

```
FROM user_usage
```

```
GROUP BY use_id
```

```
""")
```

6.06 ms  $\pm$  23.7  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 100 loops each)